

Introduction

Mercor

Project Artemis

Expert Onboarding Guide

Training frontier AI models through expert scientific reasoning

Mercor · 2026 · Confidential — authorized contractors only


How This Guide Is Structured

This guide follows your workflow from first login to first approved task. Complete the phases in order, as each one builds on the last.

Phases 0 and 1: Setup of accounts and access

Phases 2 and 3: Understanding the concepts and tasking workflow

Phase 4: The review loop, which you may cycle through to get your task approved



Ready to start on [Phase 0: Getting Started](#)→

Phase 0: Getting Started

PHASE 0

Getting Started

 After this phase, you will be able to:

- Explain what Project Artemis produces and why it matters
- Know who to contact for every type of question
- Understand your output expectations and pay model

Your job is to write STEM-based questions so difficult that a frontier AI model still fails - and to explain where and why it broke down.

What is Project Artemis?

In this project, you will help a cutting-edge AI model improve its performance in your scientific domain. You will work alongside highly qualified experts across STEM fields - Biology, Computer Science, Engineering, Physics, and Math - to craft graduate-level scientific problems drawn from open-source publications.

The way we improve the model is by finding places it is currently weak. Practically speaking, this means asking extremely challenging questions which it fails to reason through and answer correctly. When we find reasoning weaknesses within the model, we can provide correct solutions, which allow for improved training.

Meet the Team

Role	Name(s)
Strategic Project Lead	Kutay Gokcen / Luke Martocchio
Technical Program Manager (hours, contracts, Insightful)	Netra Gupta
Expert Project Managers	Vaishnavi Pai · Melissa Herman · Zach Austin · Tim Grotzer
Biology Domain Leads	Christian Stackhouse · Rachel Terry
Computer Science Domain Leads	Bailey Eccles · Yimeng Dai
Engineering Domain Lead	Mostafa Lutfi
Math Domain Leads	Aluna Rizzoli · Chen Chen · Ethan Alwaise · Ethan Robinett
Physics Domain Leads	Andre Grossardt · Michael Childress · Frank Petriello

Contact Routing

Question type	Where to go first
Most task/help questions	#artemis-help + tag @Maven, then Domain Lead or EPM if unresolved
Domain-specific questions	Your domain Slack channel, then to your Domain Lead
Access issues	#artemis-help + tag @Maven, then Domain Lead or EPM
General questions about contract/pay/hours	#artemis-help + tag @Maven, then EPM
Specific issues about contract/pay/hours	TPM - Netra Gupta

Do not message SPLs/TPMs/EPMs without first asking in the #Artemis-help channel, and tagging @Maven or messaging your Domain Lead first, but they can be contacted if you need help with:

- 💬 Contract/pay questions
- 💬 Hour increases/Insightful issues
- 💬 Things you can't get help with on #artemis-help

⚠ Confidentiality: You may never mention the name of the AI lab or any project-specific details publicly. Unauthorised disclosure or distribution of project materials is a breach of contract and will result in termination.

What is Mercor?

About Mercor

Mercor is a talent marketplace connecting companies with highly skilled contractors. For this project, Mercor connects domain experts with a leading AI lab to improve their frontier model. You can find more details about Mercor and helpful information at these links:

Key resources: Talent Portal · talent.docs.mercor.com/welcome

Insightful setup · talent.docs.mercor.com/how-to/use-insightful

Stripe payments · talent.docs.mercor.com/how-to/payments

Referral policy · talent.docs.mercor.com/policies/referrals

Tax information · talent.docs.mercor.com/policies/taxes-work-authorization

What can I tell people about my work?

- You may never mention the name of the AI Lab you work with or project specific details
- You can post on LinkedIn that you worked as a contractor for Mercor Intelligence

[Social Media Guidelines](#)

Expectations

We encourage writers to complete as many tasks as possible

- We expect each task to take roughly 4.5–5.5 hours of active work on average.
 - First tasks are likely to take longer, but time should come down with experience
- Active work time is what matters - waiting for reviewer replies or model runs does not count toward that estimate.
- All active work **must** be tracked in Insightful - this is required for all writers, reviewers, and team leads.
- New writers should wait for their first task to be approved before claiming additional tasks, and after this can work on up to 3 tasks simultaneously.

AI Use Policy

No external LLMs are permitted for any part of task creation

- You may use the LLM chatbot within RL Studio only.
- The chatbot is a useful tool for:
 - Brainstorming
 - Checking your prompt for ambiguity
 - Summarising papers
 - Grammar/spelling checks
 - Probing the difficulty of your question
- The chatbot should not be used for:
 - Fully generating your prompt
 - Fully generating your explanation
 - Finding the answer to the question

Find the Chatbot

The chatbot can be found either by clicking into any completed trajectory, or by clicking 'run batch' on the proxy run section. Clicking either #1 or #2 here will bring up the window

with the chatbot.

PROXY RUNS

Run batches of 2 quick model attempts to gauge difficulty. Mark a batch i

Run New Batch (2) **⊗ Cancel Running (2)**

Batch 1 **Running**

Golden: n,n,n,n,y,y

#1 **Running**

#2 **Running**

HOME TASK TRAJECTORY
All Tasks > Task task_022 > Trajectory traj_189

Save Changes

unknown

Task e6yaddbe / Trajectory traj_189cf43ceb224692994d13c713fb65d9 Running < Trajectory 267 / 268 >

unknown / unknown · ReAct Toolbelt Agent · code execution + filesystem + web search

Started 2m ago

Info Task Logs Files Turns Verifiers Analysis Annotation **Chat** + >

PROMPT

System
[RLS_PROBLEM_ID:task_e6yaddbe_with_search]
[RLS_POLL_TIMEOUT:0]

You are an agent that completes tasks independently.
Use the tools provided to you to complete the task to the best of your ability.

User
A three-member family of nuclear receptors (X, Y, Z) are known to be involved in retinoic signaling by heterodimerization with nuclear receptors from other receptor families. Receptor binding analysis revealed that members of this family of nuclear receptors exclusively bind to alitretinoin, but not its other geometric isomers.

These receptors have been implicated in various neurodevelopmental processes. Striatal neurons derived from adult mice with haploinsufficiency of receptor X and complete

> Show more

DIFF

Waiting for trajectory to finish

Trajectory QA Current trajectory

> Chat has access to 9 context sources

Ask about this trajectory

Ask questions about what happened during the agent's execution, timing, errors, tool usage, and more.

Type `{}` to reference variables like prompt, custom fields, or verifiers.

Drag the `||` handle on any log step into the chat to attach it as context.

Ask about this trajectory... (type "{ " for variables)

AI-generated responses may not always be accurate.

You own everything you submit. Verify every part of your response. The chatbot is used as an assistant to improve efficiency, not as a ghostwriter.

⚠ Do not assume LLM-generated content is correct - always fact-check. Repeated low-quality LLM-assisted submissions can lead to offboarding.

Pay

Pay is processed via Stripe

- All tasks approved by EOD Wednesday PT are included in that cycle, and pay should be received the following week.
- There are two rates depending on task type:
 - No tools: \$600 per approved task (standard STEM reasoning, no simulation or coding required)
 - With tools: \$700 per approved task (task requires Python, numerical simulation, or bash to solve correctly)
- In-progress tasks can be monitored on the Payments [dashboard on RLS](#).
- Reviewers and Team Leads are paid at their contracted hourly rate.

UP NEXT..... Ready to move to [Phase 1: Setup](#)→

Phase 1: Setup

PHASE 1

Onboarding & Setup



After this phase, you will be able to:

- Complete all onboarding steps in the correct order
- Access RL Studio, Slack, and Insightful without issues
- Be ready to claim and submit your first task

Quick Reference — Onboarding Steps

Complete these steps in order. You cannot submit tasks until RL Studio access is granted, which happens after the onboarding quiz.

0	Set up Insightful Install the time tracker before doing anything else. Linux users: use Xorg, not Wayland.
1	Accept your contract Sign by following the email link as soon as it arrives so you can begin working.
2	Take the Course in WorkRamp Pass at 80% (you will have 2 attempts)
3	Attend an onboarding session Required before contributing. Sessions are recurring — join at least one. Fill out the attendance tracker provided in the call.
4	Join Slack Via autoprovisioning after your session.
5	Read this guide end-to-end Also read your domain-specific instruction document.
6	Get RL Studio access Via Okta Dashboard after onboarding call. Allow up to 24 hours for access to propagate.

7

Submit your first task

Follow the Task Writing Workflow in Phase 3.

8

Adjust hours if needed

Contact TPM (Netra Gupta) on Slack.

Step 0 — Set Up Insightful

You should have received an Insightful setup email with your job offer. The app is called "Workpuls" once downloaded. Start the timer when you begin active work on any task, and stop it as soon as you stop. Do not leave it running while you are not actively working.

Writers must track time even on the pay-per-task model - this feeds project metrics. Reviewers and Team Leads must track time to receive payment.

⚠️ Fraudulent activity - logging hours without working - is grounds for immediate termination. Missed time cannot be added retroactively. Linux users: Insightful is not compatible with Wayland. Use Xorg as your display server.

Step 1 - Accept Your Contract

Accept via the email sent with your offer. Do this as soon as possible - you cannot log hours until the contract is signed.

Step 2 + 3 - Take the course and attend an Onboarding Session

You must attend a live onboarding session before contributing to the project. Sessions are held regularly - join at least one. After attending, fill out the attendance tracker link provided in the session. You will receive a follow-up email with next steps and your RL Studio access timeline.

Step 4 - Familiarize with Slack


Slack is how 99% of project communication, including major announcements, happens. Check your DMs and relevant channels daily to ensure you do not miss anything.

⚠️ Your Slack display name must match your RL Studio name so we can identify you across both platforms. On RL Studio, your name will be, e.g., "[EXP] Melissa H". Please ensure your Slack display name matches.


Channel	Purpose
#artemis-announcements	Main project announcements. DO NOT mute this channel.
#artemis-help	Primary contact point for any questions about tasks, guidelines, or issues. Always tag @Maven when asking a question to get a quick response.
#artemis-model-run-issues	Flag model or RLS issues so we can monitor system health- tag @ModelRunIssue or @RLSIssue with your task number/issue.
#artemis-general	General discussion
Domain channels (bio, cs, engineering, math, physics)	Domain-specific questions - your Domain Leads monitor these

Slack Autoprovisioning

1. Go to work.mercor.com and log in with your personal email
2. Navigate to Home → Contracts → click your Artemis contract ("Biology Expert" or similar)



3. Scroll to "Tools & Integrations" → find Okta → click "Dashboard"
4. On the Okta Dashboard, click the "Slack Enterprise" icon





Find your Mercor Expert email

Click the 3 dots next to the Google Workspace icon → email appears in the side panel. This will not be something you have created, but assigned to you, in the format word.word.word@mercortexpert.



Set your display name on Slack

Click your name in the top right corner and select Account Settings → Edit Profile → Edit Personal Information → Display Name → Save

Use the format 'Firstname lastname' or 'Firstname Last initial', e.g., 'Melissa Herman' or 'Melissa H'. Make sure your name matches the name in RL Studio!

You can also include your title, if you are a reviewer or domain lead, e.g., 'Melissa H-Biology Lead' or 'Melissa H- Biology Reviewer'

Dashboard

My Apps

Work





Add section

Notifications **1**

Add apps

My Apps

Work



Madelaine "Paige"
c-mercor



Madelaine "Paige" Thorpe
(EPM)

[Settings](#)
[Preferences](#)
[Recent Activity](#)
[Sign out](#)


Okta Personal

Dashboard

My Apps

Work

Add section

Notifications **1**

Add apps

Account

Edit Profile

Personal Information

First name
Madelaine "Paige"

Last name
Thorpe (EPM)

Okta username

Primary email

mercorUID

Display Name
Madelaine T (EPM)

Last Name Privacy
Show Full Last Name

Change your display name!
examples:
"FirstName LastName - Glacier Writer"
"FirstName LastName - Glacier Reviewer"

**if you have issues and it won't let you, you have to: change your privacy settings for last names on Okta!

Last sign in: a few seconds ago


© 2026 Okta, Inc.

Privacy

Resetting Your Okta Password

⚠ Do NOT reset your password via the Okta portal directly - this causes errors. Use the steps below only.

1. Follow the steps above to reach your Artemis contract page
2. Scroll to Tools & Integrations → Okta → click "Reset Password"



3. Set up a new passkey via Chrome using Google Password Manager

After resetting, you can authenticate into Okta and access Slack. If you see "No available workspaces," wait up to 24 hours for Slack to sync.

Steps 5 - Read the Docs

Read this entire guide and your domain-specific instruction document. Both are required before you can begin tasking - the domain doc covers notation standards, worked examples, and quality bars specific to your field.

Step 6 - Access RL Studio

Mercor RL Studio is your central platform for all project work. You write questions, run model evaluations, and submit tasks entirely within it. Access the platform via the Studio tile in your Okta Dashboard. Allow up to 24 hours after passing the quiz for access to propagate.

M Studio

Annotator

Tasks

Tasks

Worlds

Claim Worlds

Dashboards

Unclaimed Tasks

My Tasks

EPM QC (Math)

Payment Summary

Project

Artemis - Prometheus



Tasks

Create and manage tasks for annotation • 333 tasks in this world

Math

66 Unowned

Filter

Columns

Group

Page 1 of 4

333 tasks • Showing 1–100 of 333

100 rows

Task Name	Task Status	Created By	Updated By	Needs Attention From	
Task z4rzefc4	Delivered 4/13	Jordan Walsh	Subrahmanya Venkata Sai Rohit Ganti	Jordan Walsh	0
Task dbdb658d	Delivered 5/21	[EXP] Dylan J	Jordan Walsh	—	0
Task d5ppb8ea	Delivered 5/1	[EXP] Aluna R	Jordan Walsh	Jordan Walsh	0
Task dckqe5d6	Pending	[EXP] Ryan W	[EXP] Ryan W	[EXP] Ryan W	0
Task jyca592b	Pending	[EXP] Andreas Rasvanis	[EXP] Andreas Rasvanis	[EXP] Andreas Rasvanis	0
Task b96pfb71	QC Approved	[EXP] Hamid U	[EXP] Vaishnavi Pai	—	0
Task z0380463	Discarded	[EXP] Hamed H	Netra Gupta	—	0
Task qki5ee8e	In Revision	[EXP] Lukasz Plociniczak	[EXP] Melissa Herman	[EXP] Lukasz Plociniczak	0
Math Task 49	Delivered 5/11	[EXP] Armando V	Jordan Walsh	—	0
Task cl494e9d	Pending	[EXP] Constantinos Papachristoforou	[EXP] Constantinos Papachristoforou	[EXP] Constantinos Papachristoforou	0
Task kt2o2947	Pending	[EXP] Hang J	[EXP] Hang J	[EXP] Hang J	0
Task k3i3695d	Pending	[EXP] Hamid U	[EXP] Hamid U	[EXP] Hamid U	0

UP NEXT..... Ready to move to [Phase 2: Understanding the Task](#) →

Phase 2: Understanding the Task

PHASE 2

Understanding the Task

 After this phase, you will be able to:

- Articulate what makes a valid Artemis task
- Distinguish a Tier 3 reasoning failure from a Tier 2 computation error
- Understand whether your question is a reasoning challenge or a recall question
- Know the difficulty calibration target and what "In Distribution" means

The Core Task Mental Model

Read this section carefully. It gives you the mental model you need to make good judgment calls throughout the writing process, and to strengthen the quality of your task.

What you are building

An Artemis task is a scientific question derived from a specific paper, which has exactly one unambiguous, verifiable answer. It requires **reasoning** with the paper's content, not just extraction of information, and is designed so that:

- **WITHOUT the paper:** the AI model almost never solves it correctly (< 12.5% of attempts)
- **WITH the paper:** the AI model sometimes — but not always — solves it (25–87.5% of attempts)

 How to achieve this

The paper provides specialised information or details that the model can use as a **hint** or **guide** toward the correct answer.

Your question forces the model to **reason** with the content - applying, extrapolating, combining, or transforming the information to reach a new conclusion rather than using a value directly from the paper.

The Difficulty Target

You are aiming for a specific level of difficulty, verified by running three sets of AI models against your task. Every time you run a set, the model will attempt to answer your question 8 times (16 in Biology). Each set of models, as below, has access to different information and tools.

Run type	What it is	Standard target	Biology target	What it means
Baseline	No access to paper or websearch	0–1 correct out of 8 runs	0–2 correct out of 16 runs	Model rarely solves it without the paper
With paper	Direct access to PDF of paper	2–7 correct out of 8 runs	3–13 correct out of 16 runs	Model sometimes solves it with the paper, but not always
With web search	Access to websearch only	No requirement	No requirement	Informational only

Both the baseline and with-paper targets must be met **before** you can submit.



In Distribution

'In Distribution' simply means your prompt has satisfied the correct answer requirements listed in the above table. If baseline has between 0–1 correct answers out of 8, and with-paper has between 2–7 correct answers out of 8, your task is 'In Distribution'. If not, it is called 'Out of Distribution' and cannot be submitted.

⚠ If the model is getting 0/8 correct even with the paper, your task may be unclear, ambiguous, or over-stumped. Check your prompt for clarity, or add a strategic hint/simplify one reasoning step. If it is getting 8/8 correct with the paper, it is under-stumped - add complexity using one of the strategies in this section.

What Kind of Model Failure Are We Targeting?

For project Artemis, we are looking for Tier 3 failures — genuine reasoning breakdowns. Understanding the tiers helps you design questions that hit the right target.

Tier	Type	Examples	Counts for Artemis?
1	Minor Slips	Typos, phrasing errors, formatting issues	✗ No — too surface-level issues
2	Non-Reasoning Errors	Factual errors, arithmetic slips, retrieval errors, instruction-following failures	✗ No — knowledge or instruction gap, not reasoning gap
3	Major Reasoning Failures	Misconnected concepts, invalid logical leaps, multi-step chain collapse, flawed inferences	✓ Yes — this is what the difficulty distribution measures

How to recognise a Tier 3 failure

A Tier 3 failure is one where the model tries to answer the question by following multiple steps of reasoning, but has had a breakdown in reasoning along the way.

If the model simply recalled the wrong fact or made an arithmetic slip, that is Tier 1 or 2. Revise your question to require more reasoning before testing again.

Reasoning vs. Recall - The Core Distinction

This is the most important concept on the project. Simple recall questions will not be difficult enough for this project. You will need to layer in complexity via reasoning challenges, not just information extraction. If the answer to your prompt is found directly in the paper, you will need to modify it to require reasoning.

✗ Weak

RECALL: "What is the KD value of antibody A?"

Baseline model must calculate the KD from given info. The with-paper model only needs to locate the value in the paper.

Why it's weak: *The answer is directly in the paper. Having the paper is sufficient, and no reasoning required.*

Why it matters: Models with the paper will score 100% - **above** the required 25-87.5% range.

✓ **Stronger**

REASONING: "Using the binding parameters for antibody A, determine the fold-change in free antibody concentration required to shift fractional receptor occupancy from 20% to 80%."

Baseline model must calculate binding parameters **and** the fold change. The with-paper model is given initial binding parameters but must apply them to a new calculation.

Why it's stronger: Having the paper is helpful but not sufficient. The model must reason through a new calculation using paper data as a hint.

Why it matters: This means the with-paper model already has half the question worked out- an advantage over the baseline, but still must complete a reasoning step. This helps to achieve the 25-87.5% correct range.

What Counts as Reasoning: The 10 Reasoning Types

This section will detail different reasoning types you can employ in your tasks. The strategies for adding complexity follow after.

Reasoning Type	Definition	Example prompt phrasing
Deductive	Drawing specific conclusions from general laws or principles.	"Given this reaction mechanism, what is the overall reaction order?"
Inductive	Generalizing from patterns or experimental observations.	"Given these measurements, what is the value of the rate constant that fits the observed trend?"
Temporal	Predicting events or states based on order in time.	"If A is activated before B, and B represses C, by what factor does C's steady-state level change?"
Spatial	Understanding structures, orientation, or symmetry.	"Which product isomers exhibit C_{2v} symmetry?"
Causal	Identifying cause-and-effect relationships.	"What change in gene expression would cause this phenotype?"
Comparative	Judging between alternatives or evaluating differences across conditions.	"Which treatment has the highest efficacy, and what percent is it?"

Reasoning Type	Definition	Example prompt phrasing
Abstract	Working with non-concrete or theoretical ideas.	"What is the entropy change in a hypothetical closed system with constraint X?"
Pattern Recognition	Spotting and interpreting regularities in data or sequences.	"Which degenerate codons result in histidine insertion?"
Statistical	Using data, probabilities, and distributions to reach conclusions.	"What is the statistical likelihood the library diversity meets the threshold?"
Hypothetical	Predicting outcomes under hypothetical or counterfactual scenarios.	"If the methylation site were removed, by what factor would the binding affinity (Kd) change?"

Building Reasoning Complexity

This section answers the most common mid-task question: my task is too easy, what do I do? Use this section whenever the model is answering your question correctly 8/8 times in the With Paper run.



Adding Complexity

Complexity means adding more reasoning steps. Each additional step requires the model to connect concepts, apply a rule, or cross a disciplinary boundary - compounding the chance of a chain breakdown.

6 Strategies for Building Reasoning Complexity

These are the techniques to reach for when your task is not hard enough. Each one adds reasoning load without simply making the question more obscure.

1. Reframe facts into reasoning challenges

Take a fact from the paper and embed it in a scenario that requires the model to infer something new from it. Instead of "what is X," ask "given X, what follows when condition Y is applied?" The model must apply known information rather than retrieve it.

2. Add layers of reasoning

Build the complexity by introducing a second step that requires a different reasoning type than the first. If the first step is deductive, add a statistical or spatial step on top of it. Each additional reasoning type multiplies the chance of chain breakdown. When doing this, ensure that the second reasoning step requires the correct output from the first- this is how you build a more complex task.

3. Construct a realistic but hypothetical scenario

Invent a plausible situation that isn't in the literature and tie it down with specific conditions or values. The model can't recall an answer to a scenario that doesn't exist, so it has to apply domain principles from scratch rather than pattern-match a memorised one.

4. Use exceptions to rules

Build a question around a known exception, edge case, or counterexample to a general rule in your domain. Models are biased toward generalizations - they have been trained on vast amounts of text asserting the general rule. Questions that require recognising and applying an exception consistently expose this bias.

5. Cross boundaries between disciplines or scales

Require the model to connect information across two compartments it normally keeps separate. That boundary can be horizontal, such as across two subfields or methods at the same level (e.g. combining a statistical-genetics approach with a structural-biology constraint), or vertical, such as two scales of the same system (e.g. linking a systems-level phenotype down to a single molecular marker, or a bulk material property up from its microscopic origin).

6. Leverage your expertise

Write questions that only someone with your depth of training could formulate. Your domain expertise is the unfair advantage here. If an intelligent non-specialist could write or answer the same question, it is probably not hard enough. This is true particularly to questions that boil down to simple calculations or plugging information into a clear template; if it doesn't require domain-specific reasoning it will be too simple.

 If still stuck after trying 3+ strategies

Post in your domain Slack channel with the task number and a brief description of what you tried.

Domain Leads have seen hundreds of these and can usually identify the missing lever in under 5 minutes.

Iteration in Action - Watch a Prompt Get Harder

The examples below show how a simple scientific prompt is transformed into a model-stumping reasoning challenge through discrete amendments. Each row adds a named reasoning type to the previous draft.

Chemistry Example — Naphthalene Bromination to Symmetry Probability

Draft	Prompt	Reasoning added	Expected model result
1	What products form when naphthalene reacts with bromine in the presence of a Lewis acid?	Deductive	Model likely answers correctly — this is recall
2	How many monosubstitution products form, and which positions are preferentially substituted? State the ratio.	Deductive + Pattern Recognition	Model sometimes correct — borderline recall
3	Which monosubstitution product of naphthalene bromination has C_{2v} symmetry?	Spatial added	Model starts failing — spatial reasoning required
4	Given equal probability of substitution at each available position, what is the probability that random monosubstitution produces a C_{2v} symmetric product? Express as a simplified fraction.	Statistical added	Model fails frequently — multi-step chain required

Draft	Prompt	Reasoning added	Expected model result
Final	If three successive brominations occur with equal positional probability, what fraction of trisubstituted products retain at least one plane of symmetry? Express as a simplified fraction.	Comparative + Statistical	Target: 0–1 correct out of 8 baseline runs

Biology Example — Phage Display Library Diversity

Draft	Prompt	Reasoning added	Expected model result
1	What determines the size of a phage display antibody library?	Deductive	Model answers correctly — this is recall
2	A degenerate NNK primer encodes how many amino acid possibilities at each position?	Pattern Recognition	Model answers correctly — standard calculation
3	Given three degenerate NNK primers each encoding 5 positions, calculate the total theoretical sequence diversity of the combined library.	Statistical added	Model sometimes correct — multi-step calculation now required
4	If this library is expressed in 10^{10} bacteria, what fraction of the theoretical diversity is represented? Express as a percentage to 2 significant figures.	Comparative added	Model frequently fails — multi-step chain required
Final	At what minimum bacterial count does doubling the culture size no longer increase coverage by more than 5%? State as an order of magnitude.	Causal + Statistical	Target: 0–1 correct out of 8 baseline runs

UP NEXT..... Ready to move to [Phase 3: Writing your Task](#)→

Phase 3: Writing your Task

PHASE 3

Writing Your Task

 After this phase, you will be able to:

- Complete all 11 steps of the task writing workflow from claim to submission
- Produce a valid prompt, golden response and complete explanation
- Run models, interpret results, and reach In Distribution
- Use iteration strategies when your task is too easy or too hard
- Pass the pre-submission checklist before every submission

Task Writing Workflow: Quick Reference

Use the reference below to orient yourself and as a quick guide while writing. Full instructions for each step follow. There is a full video guide to the tasking workflow [here](#).

1	Claim a task Annotator → Dashboards → Unclaimed Tasks → Claim at the bottom of the task page.
2	Find your reference article arXiv, bioRxiv, or medRxiv only. Run DOI Validation + URL Uniqueness. Upload the PDF.
3	Write your task prompt Follow Tasking Guidelines below. Remember to SAVE.
4	Run Task Similarity Check Sparkles icon → Task Similarity Check. Must pass before golden response or model runs unlock.
5	Add your Golden Response Final answer only. Remember to SAVE. Models cannot run without this.
6	Run Task Analysis Click "Run Analysis." Make changes if needed.
7	Set Tool Use TRUE to run with tools, FALSE to run without.

8	Run models Run Baseline, With Paper, With Web Search. Fetch results for each. Green "In Distribution" = proceed.
9	Fill in remaining fields Task Author, Source Content Intact, Self-Containment, Explanation, Reference Location.
10	Run QA Job Report QC Report → Run → Fetch. Address or formally dispute all flags.
11	Submit for Review Click "Submit for Review" at the bottom. Check Slack for reviewer feedback.

Step 1 — Claim a Task

Navigate to: Annotator tab → Dashboards → Unclaimed Tasks.

Select a task and click "Claim" at the very bottom of the task page.

You can hold up to 3 tasks simultaneously once your first task is approved.

The screenshot shows the Studio interface. On the left sidebar, under the 'Annotator' section, 'Dashboards' and 'Unclaimed Tasks' are highlighted with red boxes. The main content area is titled 'Unclaimed Tasks' and shows a table of tasks. The first row of the table has its 'Task Status' column highlighted with a red box, showing 'Unclaimed'. A red circle highlights a 'Claim' button on the right side of the table.

Task Name	Task Status	Created By	Updated By	Needs Attention From	Criteria	Latest Score	Latest Model	GT Gr	5/7/2
Task b4vs2983	Unclaimed	[EXP] Melissa Herman	[EXP] Melissa Herman	[EXP] Melissa Herman	—	—	—	—	5/7/2
Task bexqef5	Unclaimed	[EXP] Melissa Herman	[EXP] Melissa Herman	[EXP] Melissa Herman	—	—	—	—	5/7/2
Task b75raca1	Unclaimed	[EXP] Melissa Herman	[EXP] Melissa Herman	[EXP] Melissa Herman	—	—	—	—	5/7/2
Math Task 45	Unclaimed	[EXP] Swapnil G	[EXP] Swapnil G		—	—	—	—	4/20,
Task xarfb55	Unclaimed	[EXP] Melissa Herman	[EXP] Melissa Herman	[EXP] Melissa Herman	—	—	—	—	5/9/2
Task bd095975	Unclaimed	[EXP] Melissa Herman	[EXP] Melissa Herman	[EXP] Melissa Herman	—	—	—	—	5/7/2
Task bqg5ab4a	Unclaimed	[EXP] Melissa Herman	[EXP] Melissa Herman	[EXP] Melissa Herman	—	—	—	—	5/7/2
Task xl4a6919	Unclaimed	[EXP] Melissa Herman	[EXP] Melissa Herman	[EXP] Melissa Herman	—	—	—	—	5/9/2
Task xip63a24	Unclaimed	[EXP] Melissa Herman	[EXP] Melissa Herman	[EXP] Melissa Herman	—	—	—	—	5/9/2
Task xhchb402	Unclaimed	[EXP] Melissa Herman	[EXP] Melissa Herman	[EXP] Melissa Herman	—	—	—	—	5/9/2
Task ymnm4ac3	Unclaimed	[EXP] Melissa Herman	[EXP] Melissa Herman	[EXP] Melissa Herman	—	—	—	—	5/9/2

Step 2 — Find a Reference Article

You will need first to find a source paper from arXiv, bioRxiv, or medRxiv only.

Requirements as follows:

- Each paper can only be used for one task

- Add the URL into the 'Paper Web URL' field, and click 'URL Uniqueness'
- If the automation shows 'Unique' you can continue; if it shows 'Duplicate' you will need to choose a new paper

Paper Web URL*

Note: please follow instructions carefully – this MUST be filled out in order to begin working on your task, and you may be prevented from working on a task if the paper has been already been utilized multiple times.

Provide the direct link to the pdf on the arXiv/biorXiv website by clicking "View PDF" and copying the link once the PDF is open in the browser. The link should look something like: <https://www.biorxiv.org/content/10.64898/2026.03.21.713366v1.full.pdf>

URL Uniqueness [Automation]*

- You may use up to 2 articles per task
- The PDF must be under 100 pages and may not be altered

Next, you will upload a PDF into 'Task Input Files', and click SAVE

Task Input Files

Files provided as input for this task in the filesystem

1 files, 216.1 KB


216.1 KB

Step 3 — Write Your Task Prompt

Fill in the Task Prompt field. Every prompt should have 3 components- context required to answer the question, formatting requirements for the answer, and the question itself. Ensure you follow all guidelines:

- Around 300 words (no hard maximum), free from spelling and grammar errors
- Cannot be true/false, multiple choice, or open ended ("Explain" questions not allowed)
- Question is completely unambiguous and has a single verifiable answer

- Only one defensible way to interpret the question
 - State any required assumptions or variables
- Requires reasoning rather than just data extraction
- Must be possible to answer without the paper, even if very difficult
- Contains all requirements for answer formatting- including enough detail to allow an exact string match
 - "12", "12.0", "12mg", and "the answer is 12" are NOT exact string matches
 - Desired order of the answer should be stated if more than one value is included
 - Include a statement like, 'Answer in mg as a single integer to 2 significant figures, and do not include units. Provide answers in alphabetical order, separated by a comma and space.'
- Must end in the question




The screenshot shows a web interface titled "TASK FIELDS". Under the "Task" section, there is a prompt to "Fill in task prompt and other information". Below this is a "Task Prompt" section with a text input field. The input field contains the text "New task - please update this prompt with task instructions." and is preceded by a "User" label. At the bottom left of the interface, it says "Question Word Count: 10".

Step 4 — Run Task Similarity Check

In order to see the Golden Response and model run fields, you will need to run the task similarity check.

- Click the sparkles icon in the top right and select "Task Similarity Check."
- If it fails, your prompt is too similar to an existing task in the system - revise it meaningfully and re-run before proceeding.



Task Similarity Check

Verdict: pass / warning / fail / error.

pass

Step 5 — Add Your Golden Response

Before running the models, you will need to enter your Golden Response. This is your final answer- the single correct response to your prompt.

- Scroll to Golden Response field and click 'Add Golden Response'
- Enter your final answer only and click SAVE

GOLDEN RESPONSE

Golden Response

Expected outputs that the agent should produce for this task

No golden response yet. Add one to define the expected output for this task.

+ Add Golden Response

Ensure your Golden Response fits all requirements:

- Maximum of 60 characters
- Can be a numerical answer, expression, short technical word/phrase
 - If numeric, must be an integer, reduced fraction, or radical
 - No decimals allowed
- Maximum 4 significant figures
- Does not contain any units
- Is not one of these commonly guessed answers:
 - 0 , 1 , -1 , i , π , 2π , $\pi/2$, $90^\circ/180^\circ/360^\circ$, $\pm \infty$


⚠ The golden response must be exact. Even mathematically equivalent forms fail autograding: $\frac{1}{2} \neq 0.5$, $x+y \neq y+x$, $(1,2,3) \neq (1, 2, 3)$ [spacing counts], $\boxed{1234} \neq 1234$. Specify the required form in your prompt, then produce the golden response in exactly that form.

Step 6 is currently removed- please proceed to Step 7!

Step 6 — Run Task Analysis [OPTIONAL]

The Task Analysis provides a score for your task across several dimensions before you commit to running the models. You can use this to strengthen your prompt, or disregard comments you do not agree with.

- Click "Run Analysis" in the Task Analysis field
- A report appears in approximately 2 minutes
- Read it carefully and make changes if warranted
- Read it carefully and make changes if warranted



TASK ANALYSIS

 Prompt Rubric Analysis

Run Analysis

This is an automated analysis and may not capture all issues. Always review the results and verify findings independently.

Score: **100%** Approve 9/9 checks passed

 ASSESSMENT

This is an excellent, well-crafted mathematical problem in algebraic combinatorics. The prompt is clear, rigorous, and self-contained. All mathematical definitions are provided, notation is properly formatted in LaTeX, and the expected answer format is unambiguous. The problem involves sophisticated concepts from transformation monoid theory and requires computing the ranks of semidirect products, making it appropriately challenging for frontier LLMs. The problem passes all 9 criteria with full marks.

Individual Checks (9)



CRITERION 1: Answer Format Compliance

100% >

Answer format is clearly specified as an ordered quadruple of positive integers



CRITERION 2: Single, Non-Trivial Answer

100% >

The task has a definite, checkable answer - a specific ordered quadruple



CRITERION 3: All Constants Provided

100% >

All necessary mathematical definitions and constants are provided



CRITERION 4: Prompt Unambiguity

100% >

The prompt is unambiguous with all terms precisely defined



CRITERION 5: Methodological Completeness

100% >

The mathematical framework is completely specified



CRITERION 6: LaTeX Formatting

100% >

All mathematical notation is properly formatted in LaTeX

Step 7 — Set the Tool Use

Determine if your task would require the use of Python or equivalent to solve. If yes, set the flag to TRUE. If the task is solvable through standard STEM reasoning without any coding, set it to FALSE.

Step 8 — Run Models and Verify In Distribution

In the model run section, click each of the following once in order, then click "Fetch" for each batch until all results populate:

- Run Baseline (No Paper, No Search)

- Run With Paper
- Run With Web Search

The screenshot displays three task configuration panels in the MERCOR interface. Each panel includes a 'Problem ID' field, a 'Taiga Problem ID to fetch runs from' field, and a 'Fetch' button. The panels are titled 'BASELINE (NO PAPER, NO SEARCH)', 'WITH PAPER', and 'WITH WEB SEARCH'. The 'Run Baseline (No Paper, No Search)' button is circled in red in the first panel. The 'Run With Paper' button is circled in red in the second panel. The 'Run With Web Search' button is circled in red in the third panel. The 'Fetch' buttons are circled in blue in each panel. The status message 'No external runs found. Click Fetch to retrieve them.' is displayed in each panel.

Do not hit 'run' more than once. You can click 'Fetch' as often as needed to populate all run data, even if the status still says 'Fetching'.

When the models are done, the following information is available:

- A job with time and date stamp
- All 8 individual model runs for the set (or 16 in Biology)
- The model answer, and a grade of correct or incorrect

Recall that in order to submit, your task must have 0-1 correct answers in the baseline set and 2-7 correct answers in the with-paper set (0-2 and 3-13 in Biology)

Recall that in order to submit, your task must have 0-1 correct answers in the baseline set and 2-7 correct answers in the with-paper set (0-2 and 3-13 in Biology)

WITH PAPER 16 JOBS - 128 RUNS Fetch

Problem: task_b96pfb71_with_paper

Problem ID: task_b96pfb71_with_paper

Taiga Problem ID to fetch runs from

Fetches 128 run(s). All completed.

JO	May 27 05:05 AM	May 27 04:23 AM	May 26 02:58 PM	May 26 03:38 AM	May 26 12:40 AM	May 23 07:00 PM	May 15 02:59 AM	May 14 12:00 PM	May 14 03:24 AM	May 03:12
	88% avg	100% avg	100% avg	100% avg	100% avg	100% avg	75% avg	88% avg	100% avg	

ATTEMPT 1 2 3 4 5 6 7 8 Load Transcripts

completed 100%

(9, 7, 10, 7) MODEL ANSWER

38m 28s DURATION

325,371 TOKENS

6 TURNS

May 27 05:05 AM

Model: (9, 7, 10, 7) Expected: (9, 7, 10, 7)

The platform will show "In Distribution" in green when both the baseline and with-paper targets are met. This is required before you can submit.

⚠ If you change the prompt field after running models, you must re-run ALL models before you can submit. Up to 3 error runs are acceptable if the rest of the runs are In Distribution.

Taiga Grading Override

If the autograder has marked model responses incorrectly — for example, accepting a wrong answer or rejecting a correct one — use the Taiga Grading Override.

Select "Yes" from the dropdown and explain clearly why the autograder is wrong. Only use this when your task is genuinely In Distribution based on correct grading. Do not use it to circumvent the distribution requirement.

Step 9 — Fill In Remaining Fields

Before submitting, complete the following fields:

- **Task Author:** copy-paste your RL Studio name exactly - for example, "[EXP] Melissa H"
- **Source Content Intact:** select Yes to confirm your prompt does not contradict or modify the data within the reference paper
- **Is task solvable without the paper?:** Confirm that the task can be solved, no matter how difficult, without requiring the paper
- **Explanation:** Your full step-by-step solution to the problem. This should contain all intermediate steps, and enough detail for your reviewer to follow. Point form preferred.
- **Additional context and Reference Info:** The specific page, figure, table, or section where the relevant data appears in the paper, and any other information you want your reviewer to know.


Step 10 — Run QA Job Report

Before submitting, you will need to run the QA report and respond to findings.

- Scroll to QC Report → "Run QC Report" → "Fetch QC Report" → wait for results

- If the report shows "No detailed results available," there are no issues and you can proceed.
- If a report appears, expand "Task Design Review" and "Env Linter" sections, and review each flag individually.
- Agree with a flag: click thumbs up → fix the issue → re-run the QA report

- Disagree with a flag: click thumbs down → write a brief explanation of why the flag is not valid, and send comment.



- In "Reward Hacking Detection", thumbs down all issues and paste 'This is an intentional part of task design' or something similar.

Step 11 — Submit for Review

Click "Submit for Review" at the very bottom of the task page. Once submitted, ensure you continue checking Slack for reviewer feedback and respond promptly. Always ensure you bring collaborative energy to interactions with reviewers to ensure a smooth process.

Final Answer Formatting

The platform enforces exact string match. One answer string is accepted — identical in both content and form. Mathematically equivalent answers are treated as wrong if formatted differently.

💡 Exact String Match — Non-Equivalence Examples

These pairs are NOT considered equivalent by the autograder:

`\boxed{1234}` vs 1234

`\frac{1}{2}` vs 0.5

`x+y` vs `y+x`

`(1,2,3)` vs `(1, 2, 3)` [spacing matters]

10^6 vs $1e6$

$[1,2]$ vs $(1, 2)$

Specify the exact required form in your prompt, then produce the golden response in exactly that form.

For mathematical expressions: wherever possible, avoid requesting full expressions and prefer pragmatic alternatives — numeric substitution (substitute all variables and return only the resulting number), evaluation tuples (return evaluations at distinct fixed points), or asking for a specific component such as a coefficient or numerator.

If your question does require a full expression: define all variables, parameters, functions, and domains explicitly in the prompt. Combine like terms, fully expand products, order monomials lexicographically ($x > y > z$, descending total degree), write coefficients as reduced rationals (p/q not decimals), and use no spaces inside LaTeX. State these form requirements explicitly in the prompt.

LaTeX and Mathematical Notation

Use $\dots$ for inline math and $\displaystyle \dots$ for display math. Do not use $\backslash (\backslash)$ or $\backslash [\backslash]$. Escape the percent sign as $\backslash \%$ — a bare $\%$ is a LaTeX comment. Do not define new macros. Do not number equations.



Domain-specific notation rules

LaTeX notation conventions for units, distributions, probability, linear algebra, complexity classes, and domain-specific symbols are covered in your domain instruction document.

If you have questions about notation for your subfield, post in your domain Slack channel.

UP NEXT..... Ready to move to [Phase 4: Submission and Review](#)→

Phase 4: Submission and Review

PHASE 4

Submission & Review

 After this phase, you will be able to:

- Understand the three review outcomes and what triggers each path
- Collaborate effectively with your reviewer via Slack
- Know what reviewers check at each step, including the reasoning quality gates
- Use the review process as a quality improvement loop

Working with Your Reviewer

Once you submit, a reviewer with subdomain expertise will claim your task. They will introduce themselves and provide feedback on Slack, so you will know who you are working with. The review process is collaborative. Reviewers are domain experts and partners in producing the best possible task, not gatekeepers. Engage professionally, ask questions, and discuss feedback in the most open, collaborative way possible.

The Three Review Outcomes

Path	What it means	What you do
 Approve	Task meets all guidelines. Slack notification sent.	Nothing - Task is approved, so you can move to your next task.
 Edit then Approve	Minor fixable issues (typos, grammar) that the reviewer can correct without a model re-run. Reviewer fixes directly and approves.	Nothing - Slack notification with feedback and then approval.
 Return for Revisions	Significant changes needed that may require model re-runs. Reviewer contacts you on Slack first, then leaves formal comments in RLS.	Make the requested changes, re-run models if prompt has changed, resubmit. Confirm each point is addressed in task comments.

 Path 2 vs. Path 3 — the key distinction

Path 2 (Edit then Approve): Reviewer can fix it without you re-running models. Examples: typos, minor explanation rewording.

Path 3 (Return for Revisions): the fix requires changes that need model re-runs - prompt wording, context, question structure, or golden response. Always starts with a Slack conversation before any formal return.

Pre-Submission Pre-Submission Checklist

Complete every item below before clicking "Submit for Review."

Reference Article

- ☐ Source is arXiv, bioRxiv, or medRxiv only
- ☐ DOI validation shows green tick
- ☐ URL uniqueness shows "Unique"
- ☐ PDF uploaded to Task Input Files and is under 100 pages

Prompt

- ☐ Ends in a single question
- ☐ Not multiple-choice, yes/no, or true/false
- ☐ All assumptions, variables, and constants explicitly stated
- ☐ Answer formatting specified (units, significant figures, order, spacing)
- ☐ Unambiguous - only one defensible interpretation
- ☐ Solvable without the paper
- ☐ Task Similarity Check passed

Golden Response

- ☐ 60 characters or fewer
- ☐ Only one way to state the answer
- ☐ No units included

- ☐ Not a disallowed answer (0, 1, -1, i, π , 2π , $\pi/2$, $\pm\infty$, 90/180/360 degrees)
- ☐ If numeric, is an integer, fraction, or radical only - no decimals
- ☐ Has a maximum of 4 significant figures
- ☐ Exact string matches what the prompt formatting requires

Model Runs

- ☐ Baseline: 0–1 correct out of 8 (or 0–2 for Biology)
- ☐ With Paper: 2–7 correct out of 8 (or 3–13 for Biology)
- ☐ Task shows "In Distribution" in green

Explanation

- ☐ All intermediate steps shown
- ☐ Reviewer can follow the reasoning and arrive at the same answer independently
- ☐ LaTeX used correctly for all symbols and equations
- ☐ No rounding of intermediate steps

Supporting Fields

- ☐ Task Author field matches RLS name exactly
- ☐ "Source Content Intact" is selected
- ☐ "Answerable without paper" is confirmed
- ☐ "Reference Information Location" is specific
- ☐ Tool Use flag correctly set (TRUE or FALSE)
- ☐ QA Job Report run — all flags addressed or disputed with documented reasoning

Reviewer Workflow

Reviews should take under 90 minutes. The five steps below form the complete review sequence.

1	Retrieve article and initial scan Download and open the reference article. If the link is broken, the question is incomprehensible, or major access errors exist → flag, return to writer with feedback, move to next task.
2	Evaluate the prompt Clear and unambiguous. Single question. Ends in a question. Not multiple-choice or yes/no. Formatting requirements (rounding, sig figs, order) explicitly stated in the prompt.
3	Check the golden response Matches the format the prompt requires. Exact string. Not a disallowed value. No units. 60 characters or fewer. Integer, fraction, or radical only if numerical.
4	Review explanation and reasoning chain Expected failure is Tier 3 (reasoning breakdown), not Tier 2 (computation slip or knowledge gap). LaTeX compiles in Overleaf. No rounding of intermediate steps. Final answer is supported by explanation.
5	Review supporting fields All required fields filled. Self-contained. Model runs In Distribution. Reference Information Location specific enough to verify the answer.

⚠ The Gut Check: Before approving, ask: "Does this question require the model to REASON, or just locate an obscure fact?" Obscure paper content is acceptable and encouraged, but the question must ALSO require reasoning even if that information is given directly. Valid: uses rare data from paper AND requires the model to reason with it. Invalid: only tests whether the model can find or recall a rare fact.

Full Reviewer Checklist

Step 2 — Prompt

- ☐ Clear and unambiguous — only one defensible interpretation
- ☐ Formatting requirements explicitly stated (rounding, sig figs, order)
- ☐ Single question — not multiple sub-questions

- ☐ Not multiple-choice or yes/no format
- ☐ Ends in a question

Step 3 — Golden Response

- ☐ Matches the format required by the prompt exactly
- ☐ Is an exact string that can be auto-graded
- ☐ Not a disallowed value (0, 1, -1, i, π , 2π , $\pi/2$, 90° , 180° , 360° , $\pm\infty$)
- ☐ Does not contain units
- ☐ 60 characters or fewer

Step 4 — Reasoning and Explanation

- ☐ Expected model failure is Tier 3 — not a computation slip or knowledge gap
- ☐ Question requires the model to reason across data or evidence — not just extract a fact
- ☐ LaTeX compiles correctly (verify in Overleaf)
- ☐ Grammar, spelling, and capitalisation are correct
- ☐ Problem cannot be found in existing public datasets or easily searchable online
- ☐ Final answer is single, unambiguous, and definitive
- ☐ No rounding of intermediate steps — full precision throughout

Step 5 — Supporting Fields

- ☐ All required fields are filled out
- ☐ Prompt is solvable without reading the full paper
- ☐ Model runs are In Distribution
- ☐ Reference Information Location is specific (page, figure, table, section)
- ☐ Task author field matches RLS name exactly

Reference

Consult this section when you have a specific question. You do not need to read it before your first task.

Frequently Asked Questions

Onboarding

Question	Answer
Can I use the quiz to learn the guidelines, even if I fail initially?	Yes — retake as many times as needed. It is open-note.
I do not have access to Slack channels. Who do I contact?	DM Vaishnavi Pai (EPM) on Slack, or post in #artemis-help if you have any access issues.
Should I track time during onboarding?	No need to track time during onboarding sessions, reading materials, and quiz time. Track all task writing and review time.

Claiming and Managing Tasks

Question	Answer
My first task was approved. Can I claim more?	Yes - up to 3 tasks simultaneously.
Can CS writers with Electronics Engineering backgrounds claim Engineering tasks?	Yes.
Does the platform support image uploads for multimodal tasks?	No - PDFs only, not standalone images.
What settings should I use for Agent, Platform, and Bundle in Trajectory?	Leave all blank- you do not need to interact with this section.
Why can I not see the Submit button?	All required fields must be filled. Check every section of the task page.

Model Runs

Question	Answer
The model is getting 8/8 correct with the paper. What do I do?	Add complexity — use one of the 8 strategies above. Avoid simply making the question more obscure without adding a reasoning step.
The model is getting 0/8 correct even with the paper. What do I do?	Add small strategic hints until you reach the sweet spot. Check that your prompt is not unclear or ambiguous- this is often the reason for 0/8 with paper. You can also inspect the transcripts of the models to see where they are going wrong.
A few model runs (3 max) give error messages, but I would otherwise be In Distribution. Is this okay?	Yes — up to 3 error runs are acceptable.
Run With Paper models keep failing. What should I try?	Re-upload the PDF. Review your prompt and golden response for formatting issues. If unresolved, post in #artemis-model-run-issues with your task number.
Can I cancel a model run?	Not currently. Modify your task and re-run after the current run completes.

Review Process

Question	Answer
How long does review take?	Check #artemis-announcements for current turnaround estimates — timing varies across the project lifecycle.
My task was returned. Does that affect my record?	We do use revision cycles as one of our metrics, but this is a normal part of the review loop.

Common Mistakes

These are the most frequent errors that prevent first-pass approval. Review before every submission.

Most Common Rejection Reasons

- The prompt modifies or contradicts content in the reference article
- Questions are not self-contained (requires the paper to answer)
- Access to the paper does not significantly simplify the problem
- Questions rely on hidden or implicit context from the paper
- Relevant numerical values not stated in the prompt — model expected to extract from paper figures/tables
- Output format not specified (units, sig figs, rounding) → non-unique answers
- Explanation does not show calculations and reasoning steps
- Intermediate steps are rounded (carry full precision throughout)
- LaTeX guidelines not followed correctly
- Task requires Python / tools but the Tool Use flag is set to FALSE